




Objective reward (higher is better)



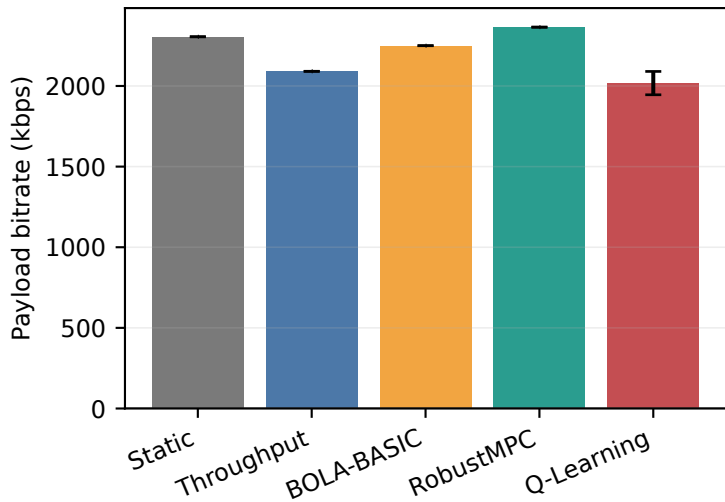
Rebuffering (lower is better)



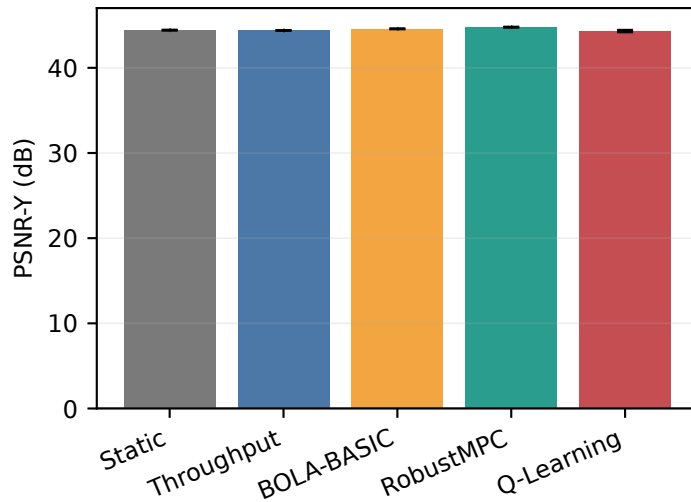
Startup delay (lower is better)



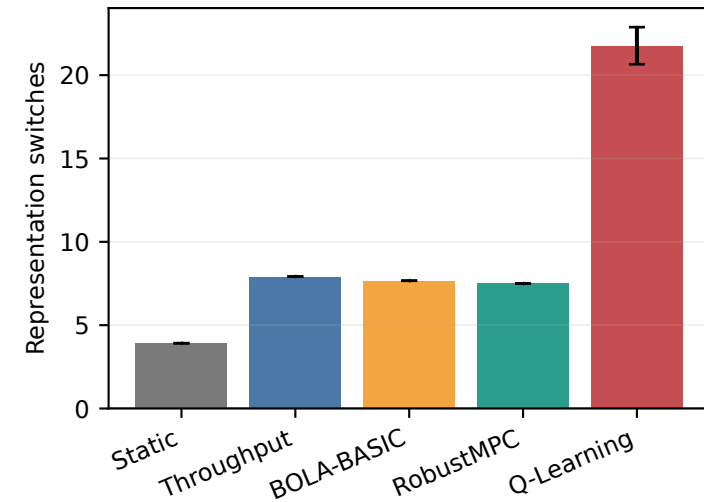
Payload bitrate (higher is better)



PSNR-Y (higher is better)



Representation switches (lower is better)



95% intervals are computed across training seeds; deterministic baselines repeat identically.